


The Story of Correlation

A Data Scientist's Guide to Measuring Relationships



Our Journey: A Quest to Measure Relationships

This presentation follows a logical progression, starting with a basic tool and building towards more sophisticated methods. Each step addresses the limitations of the previous one.

- **Part I: The First Attempt (Covariance)**

Can we measure if two variables move together?

- **Part II: The Standardized Solution (Pearson Correlation)**

How can we measure the strength of the relationship on a universal scale?

- **Part III: The Robust Alternative (Spearman Correlation)**

What if the relationship isn't a straight line?

- **Part IV: The Foundational Principle (Central Limit Theorem)**


A core statistical truth that empowers our analysis.

The Core Question: How Do Study Hours Affect Exam Scores?

Let's explore a classic data science problem. We have data for five students and want to determine the relationship between the hours they studied and their final exam scores.

- Our goal is to quantify this relationship. Does more studying lead to higher scores? By how much? How confident can we be in this conclusion?

Student	Study Hours (X)	Exam Score (Y)
1	2	65
2	4	70
3	6	78
4	8	80
5	10	85



Part I: Covariance - A Measure of Joint Variability

Covariance is our first tool. It measures the direction of the relationship between two variables.

- * **Positive Covariance:** Indicates that as one variable increases, the other tends to increase. (e.g., Study Hours and Exam Scores)
- * **Negative Covariance:** Indicates that as one variable increases, the other tends to decrease.
- * **Zero Covariance:** Indicates no linear relationship.

The distance of an individual X value from its average.

The distance of an individual Y value from its average.

$$\text{Cov}(X, Y) = \frac{1}{n - 1} \sum (X_i - \bar{X})(Y_i - \bar{Y})$$

Sum: We do this for every data point and add them all up.

Sample Size: We divide by the number of data points minus one.

Covariance in Action: Breaking Down the Calculation

Let's apply the formula to our 'Study Hours vs. Scores' data. The goal is to see how each data point deviates from the average for both variables simultaneously.

Mean Study Hours ($\bar{X}$): 6
Mean Exam Score ($\bar{Y}$): 75.6

X	Y	X - $\bar{X}$	Y - $\bar{Y}$	(X - $\bar{X}$)(Y - $\bar{Y}$)
2	65	-4	-10.6	42.4
4	70	-2	-5.6	11.2
6	78	0	2.4	0
8	80	2	4.4	8.8
10	85	4	9.4	37.6
Sum				$\Sigma = 100$

$$\text{Cov}(X, Y) = \frac{\sum [(X_i - \bar{X})(Y_i - \bar{Y})]}{n - 1}$$

$$\text{Cov}(X, Y) = \frac{100}{5 - 1} = 25$$

The Result: A Positive Relationship, But a Vague Signal

Interpretation

Our covariance is **+25**. The positive sign correctly tells us that as study hours increase, exam scores tend to increase. This confirms the direction of the relationship.

+25

The Limitation

But what does '25' actually mean? Is that strong? Weak? We can't tell.


- **Scale Dependent:** The value of covariance is dependent on the units of the variables. If we measured scores out of 1000 instead of 100, the covariance would be much larger.
- **No Standardized Bounds:** The value is not confined to a specific range. It could be 25, 1000, or 25,000. This makes it impossible to compare relationships.
- *"We only know if it's positive or negative, not how strong the relationship is."*



Part II: Pearson Correlation - Standardizing the Signal

To solve the ambiguity of covariance, we use Pearson Correlation. It takes the covariance value and standardizes it, forcing the output into a consistent and interpretable range.

The Key Upgrade: The Pearson Correlation Coefficient (denoted as r) always falls between **-1 and +1**. This gives us a universal scale to measure the **strength** and direction of a **linear** relationship.



$$r = \frac{Cov(X, Y)}{\sigma_X \cdot \sigma_Y}$$

The covariance we just calculated (**25**).


Standard Deviation: This measures the "spread" of each variable and is used to normalize the result.

Pearson in Action: From an Ambiguous 25 to a Clear 0.98

Using our calculated covariance and the standard deviations of our data, we can find the Pearson coefficient.

Step 1: Calculate Standard Deviations

- Standard Deviation of Study Hours (σ_X): **3.16**
- Standard Deviation of Exam Scores (σ_Y): **7.91**



Step 2: Apply the Pearson Formula

$$r = \frac{\text{Cov}(X, Y)}{\sigma_X * \sigma_Y}$$
$$r = \frac{25}{3.16 * 7.91}$$
$$r = \frac{25}{24.99}$$


$$r \approx 0.98$$

A value of 0.98 is very close to +1, indicating an **extremely strong positive linear relationship** between study hours and exam scores. We can now say with high confidence that more studying is associated with higher scores in this dataset.

Interpreting the Pearson Coefficient: A Visual Dictionary


The power of the Pearson r is its clear, standardized scale.

**Strong Positive
Correlation ($r \approx +1$)**




As X increases, Y consistently increases. The points form a tight line. Our result of $r = \mathbf{0.98}$ looks like this.

**Strong Negative
Correlation ($r \approx -1$)**



As X increases, Y consistently decreases. The points form a tight, downward-sloping line.


**No Correlation
($r \approx 0$)**



There is no discernible linear pattern between X and Y. The points are scattered randomly.

The Limitation of Pearson: It Only Sees Straight Lines

Pearson Correlation is excellent at measuring **linear** relationships. But what happens when the relationship is clear, but not a straight line? Pearson may underestimate the strength of a **non-linear** relationship.



"Pearson correlation may not be able to capture non-linear properties."

Part III: Spearman Correlation - Measuring Monotonic Relationships

Spearman Rank Correlation is the tool we need for non-linear relationships. It answers a slightly different question:

- **Pearson:** Do the variables move together in a straight line?
- **Spearman:** When one variable increases, does the other consistently increase (or decrease), regardless of the shape?

The Key Idea

Spearman doesn't use the actual data values. Instead, it converts all values to **ranks** and then calculates the Pearson correlation on those ranks.

Situation	Use Pearson?	Use Spearman?
Linear Relationship	✔ Yes	✔ Yes
Non-Linear Monotonic	✘ No (Inaccurate)	✔ Yes
Data with Outliers	✘ No (Sensitive)	✔ Yes
Ordinal Data	✘ No	✔ Yes


How Spearman Works: The Transformation to Ranks

The process is simple: first, we ignore the actual values and assign a rank to each data point within its own variable, from lowest to highest.

Original Data

Study Hours (X)	Exam Score (Y)
2	65
4	70
6	78
8	80
10	85

TRANSFORM
TO RANKS



Ranked Data


Rank of Hours (RX)	Rank of Score (RY)
1 (lowest)	1 (lowest)
2	2
3	3
4	4
5 (highest)	5 (highest)

Spearman then calculates the Pearson Correlation Coefficient on the **RX** and **RY** columns. For this perfectly linear data, the result is still 1.0. The real power is seen in non-linear data.

Pearson vs. Spearman: A Visual Showdown

The choice of method matters. Let's see how both perform on two different datasets.

Linear Data




Results

Pearson Result: $r = 1.0$ ✓

Spearman Result: $\rho = 1.0$ ✓

Both methods work perfectly for linear data.

Non-Linear Monotonic Data



Results

Pearson Result: $r = 0.88$ ✗

Spearman Result: $\rho = 1.0$ ✓

Only Spearman accurately captures the perfect strength of a non-linear monotonic relationship.

Part IV: The Central Limit Theorem (CLT)

Beneath many data science techniques lies a powerful statistical principle: the Central Limit Theorem.

The Theorem Stated Simply

No matter the original distribution of a population (whether it's bell-shaped, skewed, or uniform), the **distribution of the means of its samples** will be approximately a normal (bell-shaped) distribution, as long as the sample size is large enough (typically $n \geq 30$).

Key Ingredients

- Start with *any* population distribution.
- Take many random samples of a sufficient size ($n \geq 30$).
- Calculate the mean of each sample.
- Plot a histogram of all those sample means.

The Result

The histogram of sample means will form a normal distribution.

Why CLT is a Data Scientist's Superpower

The Central Limit Theorem isn't just a theory; it's what makes much of inferential statistics possible.

- **Foundation for Hypothesis Testing:** It allows us to use tests that assume normality (like t-tests) even when our original data is not normal.
- **Enables Inference:** It lets us use the mean of a single large sample to make confident inferences about the entire population's mean.
- **Reliable Predictions:** By knowing the sampling distribution is normal, we can calculate confidence intervals with a known margin of error.

